

```
In [1]: #Load required libraries

import pandas as pd
import numpy as np
import plotly.express as px
import seaborn as sns
import matplotlib.pyplot as plt
from matplotlib.ticker import NullFormatter
import opendatasets as od

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.feature_selection import chi2, SelectKBest
from sklearn.metrics import accuracy_score, roc_curve, roc_auc_score, confusion_matrix, c
from imblearn.over_sampling import SMOTE

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.naive_bayes import MultinomialNB
from sklearn.svm import SVC

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
```

```
In [2]: #Download the creditcard fraud dataset from Kaggle
#od.download("https://www.kaggle.com/datasets/fedesoriano/heart-failure-prediction/data?
#od.download("https://www.kaggle.com/code/fahadmehfoooz/heartattack-prediction-with-91-8
```

```
In [3]: ## Load the dataset
```

```
In [4]: heart_df = pd.read_csv("heart-2.csv")
heart_df.head(5)
```

```
Out[4]:
```

|   | Age | Sex | ChestPainType | RestingBP | Cholesterol | FastingBS | RestingECG | MaxHR | ExerciseAngina | Oldpeak | S |
|---|-----|-----|---------------|-----------|-------------|-----------|------------|-------|----------------|---------|---|
| 0 | 40  | M   | ATA           | 140       | 289         | 0         | Normal     | 172   | N              | 0.0     |   |
| 1 | 49  | F   | NAP           | 160       | 180         | 0         | Normal     | 156   | N              | 1.0     |   |
| 2 | 37  | M   | ATA           | 130       | 283         | 0         | ST         | 98    | N              | 0.0     |   |
| 3 | 48  | F   | ASY           | 138       | 214         | 0         | Normal     | 108   | Y              | 1.5     |   |
| 4 | 54  | M   | NAP           | 150       | 195         | 0         | Normal     | 122   | N              | 0.0     |   |

## Data Processing

### Check for null columns

```
In [5]: heart_df.isna().sum()
```

```
Out[5]:
```

|               |   |
|---------------|---|
| Age           | 0 |
| Sex           | 0 |
| ChestPainType | 0 |
| RestingBP     | 0 |
| Cholesterol   | 0 |
| FastingBS     | 0 |
| RestingECG    | 0 |
| MaxHR         | 0 |

```
ExerciseAngina    0
Oldpeak           0
ST_Slope          0
HeartDisease      0
dtype: int64
```

Chest Pain types Value 1: typical angina Value 2: atypical angina Value 3: non-anginal pain Value 4: asymptomatic

## Check for duplicates

```
In [6]: heart_df[heart_df.duplicated()==True]
```

```
Out[6]:
```

|  | Age | Sex | ChestPainType | RestingBP | Cholesterol | FastingBS | RestingECG | MaxHR | ExerciseAngina | Oldpeak | ST |
|--|-----|-----|---------------|-----------|-------------|-----------|------------|-------|----------------|---------|----|
|--|-----|-----|---------------|-----------|-------------|-----------|------------|-------|----------------|---------|----|

```
In [7]: print('Dataframe before dropping duplicates :', heart_df.shape)
heart_df = heart_df.drop_duplicates()
print('Dataframe before dropping duplicates :', heart_df.shape)
```

```
Dataframe before dropping duplicates : (918, 12)
Dataframe before dropping duplicates : (918, 12)
```

There are no duplicates to drop.

## Renaming columns

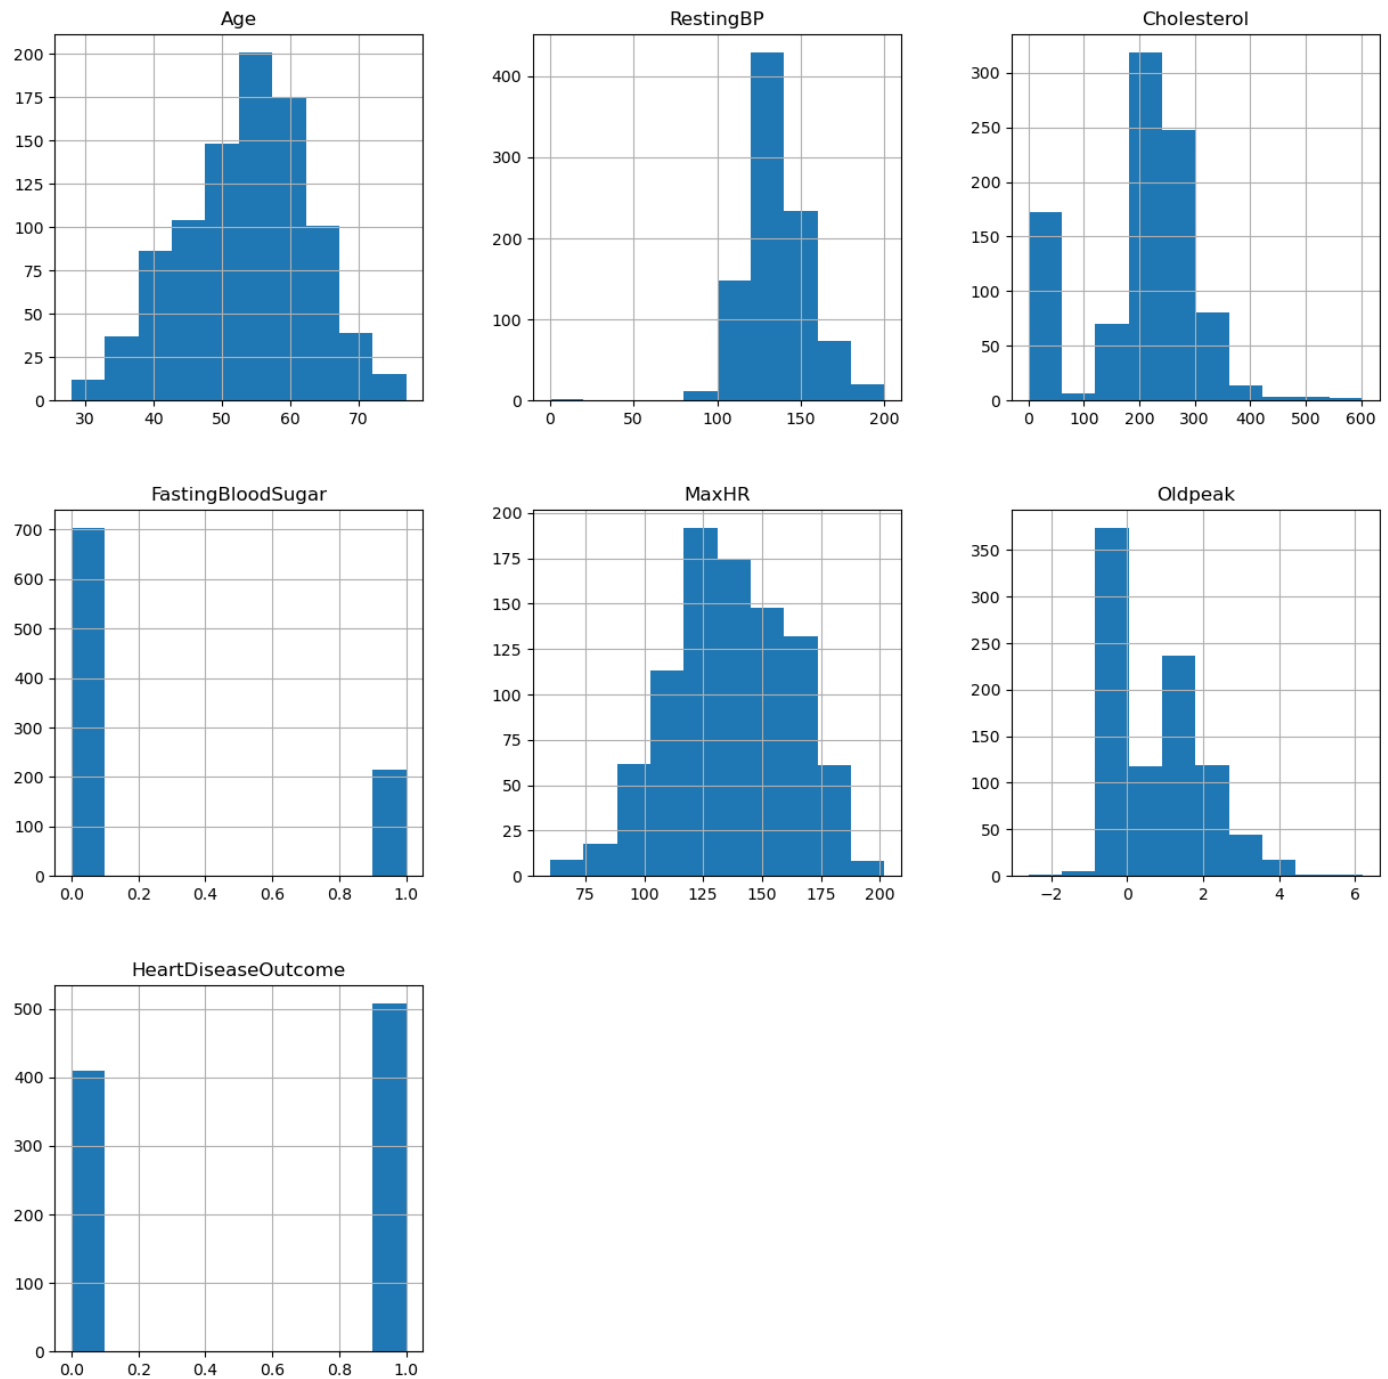
```
In [8]: heart_df.columns = ['Age', 'Sex', 'ChestPainType', 'RestingBP', 'Cholesterol', 'FastingBloodS',
                             'ST_Slope', 'HeartDiseaseOutcome']
heart_df.head(5)
```

```
Out[8]:
```

|   | Age | Sex | ChestPainType | RestingBP | Cholesterol | FastingBloodSugar | RestingECG | MaxHR | ExerciseAngina | Oldpeak | ST |
|---|-----|-----|---------------|-----------|-------------|-------------------|------------|-------|----------------|---------|----|
| 0 | 40  | M   | ATA           | 140       | 289         | 0                 | Normal     | 172   |                | N       |    |
| 1 | 49  | F   | NAP           | 160       | 180         | 0                 | Normal     | 156   |                | N       |    |
| 2 | 37  | M   | ATA           | 130       | 283         | 0                 | ST         | 98    |                | N       |    |
| 3 | 48  | F   | ASY           | 138       | 214         | 0                 | Normal     | 108   |                | Y       |    |
| 4 | 54  | M   | NAP           | 150       | 195         | 0                 | Normal     | 122   |                | N       |    |


## Data Visualizations


```
In [9]: heart_df.hist(figsize = (15,15))
plt.show()
```



```
In [10]: fig = px.pie(heart_df[heart_df.HeartDiseaseOutcome==1], names='ChestPainType', title='<
fig.update_traces(textposition='inside', textinfo='percent+label')
fig.update_layout(title = "Percentage of Heart Attacks by Chest Pain Type")
fig.show("notebook")
```

Percentage of Heart Attacks by Chest Pain Type

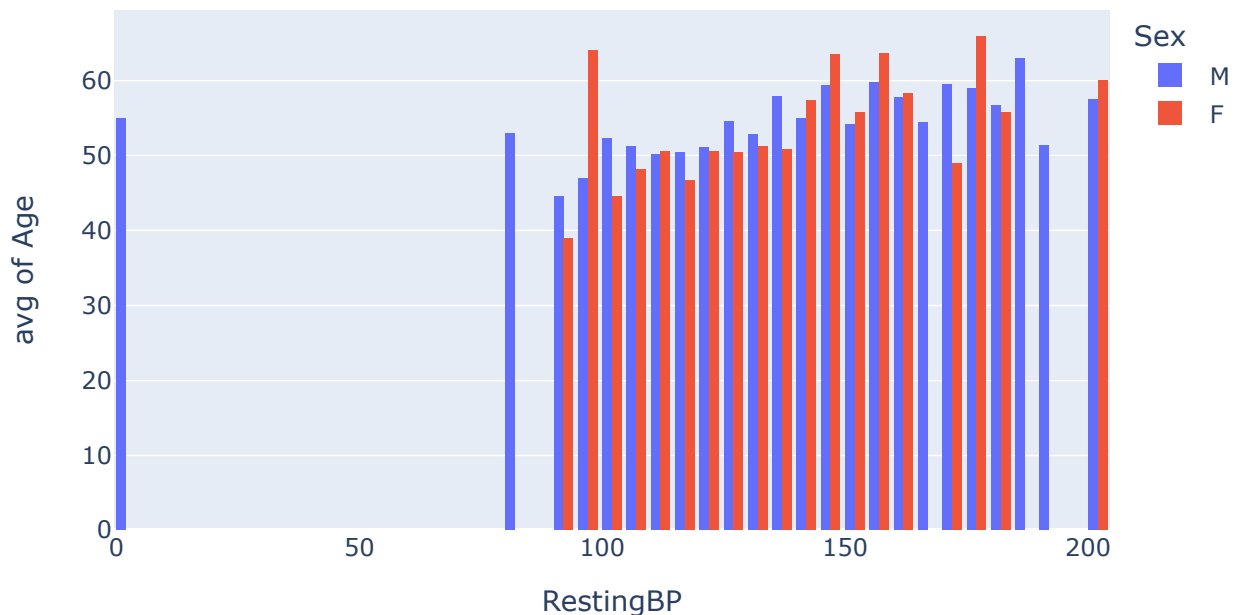




Sex: displays the gender of the individual using the following format : 1 = male 0 = female

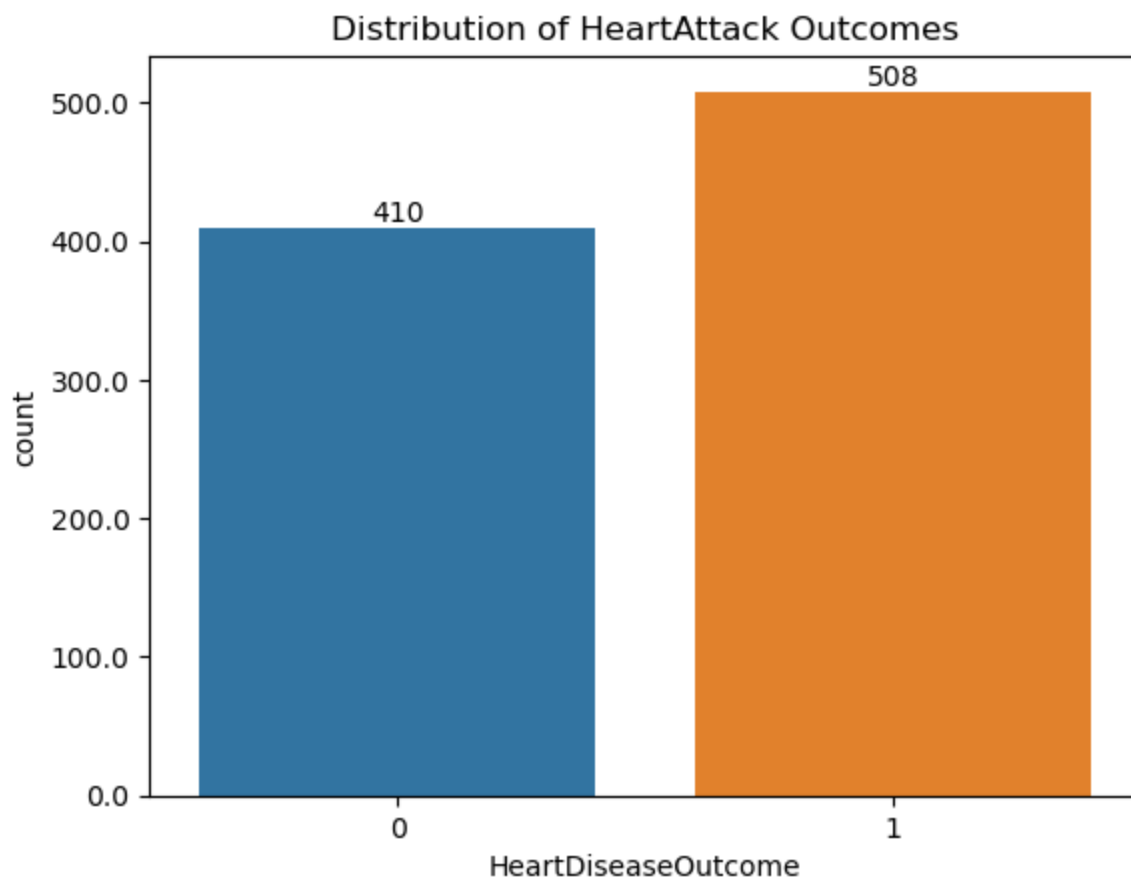
```
In [45]: fig = px.histogram(heart_df, y="Age", x="RestingBP", color='Sex', barmode='group', histfun
fig.update_layout(title = "Distribution of Resting Blood Pressure by Gender")
fig.show("notebook")
```

Distribution of Resting Blood Pressure by Gender



```
In [12]: xx = heart_df['HeartDiseaseOutcome'].value_counts().reset_index()
def formatter(x, pos):
    return str(x)

ax = sns.barplot(x="HeartDiseaseOutcome", y="count", data=xx)
ax.set_title('Distribution of HeartAttack Outcomes')
ax.yaxis.set_major_formatter(formatter)
ax.yaxis.set_minor_formatter(NullFormatter())
for i in ax.containers:
    ax.bar_label(i,)
```



## Model Building

```
In [13]: X = heart_df.drop(['HeartDiseaseOutcome'],axis=1)
Y = heart_df['HeartDiseaseOutcome']
X.shape, Y.shape
```

```
Out[13]: ((918, 11), (918,))
```

```
In [14]: X.dtypes
```

```
Out[14]: Age                int64
Sex                object
ChestPainType      object
RestingBP          int64
Cholesterol         int64
FastingBloodSugar  int64
RestingECG         object
MaxHR              int64
ExerciseAngina     object
Oldpeak            float64
ST_Slope           object
dtype: object
```

```
In [15]: # Encode categorical variables (e.g., 'gender', 'category', 'state', etc.)
categorical_columns = ['Sex', 'ChestPainType', 'RestingECG', 'ExerciseAngina', 'ST_Slope']
for col in categorical_columns:
    le = LabelEncoder()
    X[col] = le.fit_transform(X[col])
```

### Split dataset into Train and Test Sets

```
In [16]: scaler = StandardScaler()
x = scaler.fit_transform(X)
```

```
X.shape, x.shape
```

```
Out[16]: ((918, 11), (918, 11))
```

```
In [17]: # Split the data into training and testing sets
X_train, X_test, Y_train, Y_test = train_test_split(x,Y, test_size=0.2, random_state=42)
```

```
In [18]: X_train.shape, X_test.shape, Y_train.shape, Y_test.shape
```

```
Out[18]: ((734, 11), (184, 11), (734,), (184,))
```

## Models

### Random Forest

```
In [19]: # Use the RandomForestClassifier to fit balanced data
rfc = RandomForestClassifier()
rfc_model = rfc.fit(X_train,Y_train)

#Predict y data with classifier:
y_pred_rfc = rfc_model.predict(X_test)

# Evaluate the model
print(classification_report(Y_test, y_pred_rfc))
print(confusion_matrix(Y_test, y_pred_rfc))
print(f'ROC-AUC score : {roc_auc_score(Y_test, y_pred_rfc)}')
print(f'Accuracy score : {accuracy_score(Y_test, y_pred_rfc)}')
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.84      | 0.90   | 0.87     | 77      |
| 1            | 0.92      | 0.88   | 0.90     | 107     |
| accuracy     |           |        | 0.89     | 184     |
| macro avg    | 0.88      | 0.89   | 0.88     | 184     |
| weighted avg | 0.89      | 0.89   | 0.89     | 184     |

```
[[69  8]
 [13 94]]
ROC-AUC score : 0.8873042845005461
Accuracy score : 0.8858695652173914
```

```
In [20]: #Build the confusion matrix
matrix = confusion_matrix(Y_test, y_pred_rfc, labels=[1,0])


print(matrix)

# Create pandas dataframe
df = pd.DataFrame(matrix)

# Create a heatmap
sns.heatmap(df, annot=True, cbar=None, cmap="Blues",fmt='.0f')
plt.title("RandomForestClassifier Confusion Matrix"), plt.tight_layout()
plt.ylabel("True Class"), plt.xlabel("Predicted Class")
plt.show()
```

```
[[94 13]
 [ 8 69]]
```

RandomForestClassifier Confusion Matrix



## Logistic Regression

```
In [21]: # Train a logistic regression model
logistic_model = LogisticRegression(solver='liblinear', random_state=42)
logistic_model.fit(X_train,Y_train)

# Make predictions on the test set
y_pred_lr = logistic_model.predict(X_test)

# Evaluate the model
print(classification_report(Y_test, y_pred_lr))
print(confusion_matrix(Y_test, y_pred_lr))
print(f'ROC-AUC score : {roc_auc_score(Y_test, y_pred_lr)}')
print(f'Accuracy score : {accuracy_score(Y_test, y_pred_lr)}')
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.77      | 0.88   | 0.82     | 77      |
| 1            | 0.91      | 0.81   | 0.86     | 107     |
| accuracy     |           |        | 0.84     | 184     |
| macro avg    | 0.84      | 0.85   | 0.84     | 184     |
| weighted avg | 0.85      | 0.84   | 0.84     | 184     |

```
[[68  9]
 [20 87]]
ROC-AUC score : 0.8481004976332078
Accuracy score : 0.842391304347826
```


```
In [22]: #Build the confusion matrix
matrix = confusion_matrix(Y_test, y_pred_lr, labels=[1,0])

print(matrix)
```

```
# Create pandas dataframe
df = pd.DataFrame(matrix)

# Create a heatmap
sns.heatmap(df, annot=True, cbar=None, cmap="Blues",fmt='.0f')
plt.title("LogisticRegression Confusion Matrix"), plt.tight_layout()
plt.ylabel("True Class"), plt.xlabel("Predicted Class")
plt.show()
```

```
[[87 20]
 [ 9 68]]
```



## Support Vector Machine (SVM)

```
In [23]: svc_model = SVC()
svc_model.fit(X_train, Y_train)

y_pred_svc = svc_model.predict(X_test)

# Evaluate the model
print(classification_report(Y_test, y_pred_svc))
print(confusion_matrix(Y_test, y_pred_svc))
print(f'ROC-AUC score : {roc_auc_score(Y_test, y_pred_svc)}')
print(f'Accuracy score : {accuracy_score(Y_test, y_pred_svc)}')
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.82      | 0.86   | 0.84     | 77      |
| 1            | 0.89      | 0.87   | 0.88     | 107     |
| accuracy     |           |        | 0.86     | 184     |
| macro avg    | 0.86      | 0.86   | 0.86     | 184     |
| weighted avg | 0.87      | 0.86   | 0.86     | 184     |



```
[[66 11]
 [14 93]]
ROC-AUC score : 0.863150867823765
Accuracy score : 0.8641304347826086
```


```
In [24]: #Build the confusion matrix
matrix = confusion_matrix(Y_test, y_pred_svc, labels=[1,0])

print(matrix)

# Create pandas dataframe
df = pd.DataFrame(matrix)

# Create a heatmap
sns.heatmap(df, annot=True, cbar=None, cmap="Blues",fmt='.0f')
plt.title("Support Vector Machine Confusion Matrix"), plt.tight_layout()
plt.ylabel("True Class"), plt.xlabel("Predicted Class")
plt.show()
```

```
[[93 14]
 [11 66]]
```



## Naive Bayes

```
In [25]: # Initialize and train the Multinomial Naive Bayes classifier

# Ensure non-negative values in the feature vectors
x_train = np.maximum(0, X_train)
x_test = np.maximum(0, X_test)

nb_model = MultinomialNB()
nb_model.fit(x_train, Y_train)
```

```

# Make predictions on the test data
y_pred_nb = nb_model.predict(x_test)

# Evaluate the model
print(classification_report(Y_test, y_pred_nb))
print(confusion_matrix(Y_test, y_pred_nb))
print(f'ROC-AUC score : {roc_auc_score(Y_test, y_pred_nb)}')
print(f'Accuracy score : {accuracy_score(Y_test, y_pred_nb)}')

```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.73      | 0.87   | 0.79     | 77      |
| 1            | 0.89      | 0.77   | 0.82     | 107     |
| accuracy     |           |        | 0.81     | 184     |
| macro avg    | 0.81      | 0.82   | 0.81     | 184     |
| weighted avg | 0.82      | 0.81   | 0.81     | 184     |

```

[[67 10]
 [25 82]]
ROC-AUC score : 0.8182425051583929
Accuracy score : 0.8097826086956522

```

```

In [26]: #Build the confusion matrix
matrix = confusion_matrix(Y_test, y_pred_nb, labels=[1,0])

print(matrix)

# Create pandas dataframe
df = pd.DataFrame(matrix)

# Create a heatmap
sns.heatmap(df, annot=True, cbar=None, cmap="Blues", fmt='.0f')
plt.title("Naive Bayes Confusion Matrix"), plt.tight_layout()
plt.ylabel("True Class"), plt.xlabel("Predicted Class")
plt.show()

[[82 25]
 [10 67]]

```

Naive Bayes Confusion Matrix

